

# Enhancing Fault Tolerance in TMR Soft RISC-V Linux SoCs through Single-point Failure Analysis PhD Defense

Andrew E Wilson

Department of Electrical and Computer Engineering  
Brigham Young University, Provo, UT



This work has been supported by the I/UCRC program of the NSF under grant number 1738550, the Utah NASA Space Grant Consortium (UNSGC) and Los Alamos Neutron Science Center (LANSCE) under proposal numbers NS-2017-7574-F, NS-2017-7574-A and NS-2018-7895-A



# FPGAs and Soft Processors in Space

## FPGAs:

- **Custom Sensor and Actuator Integration**  
Serve as reconfigurable interfaces for mission-specific sensors and actuators.
- **Onboard Processing for Performance Optimization**  
Enable low-latency, power-efficient, and high-throughput computation tailored to mission needs.
- **Support for On-Orbit Reconfiguration**  
Allow hardware logic updates post-deployment for mission adaptability and fault recovery.



Perseverance  
Mars Lander



SpaceCube 2.0



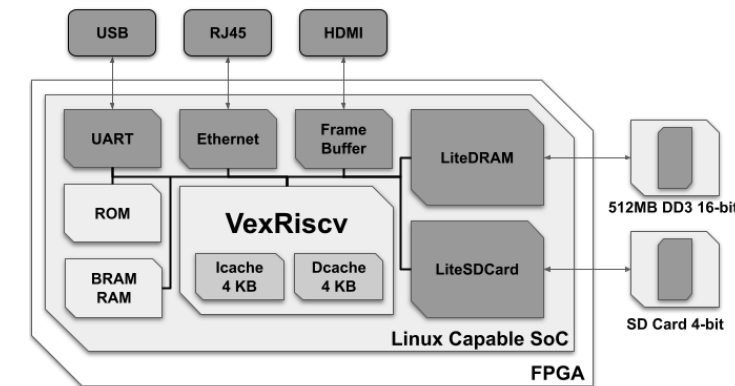
Ingenuity  
Mars Helicopter

## Soft Processors:

- **Tightly-Coupled Control Systems**  
Enable precise coordination with adjacent hardware components and real-time interfaces.
- **Software Ecosystem Compatibility**  
Support standard operating systems, libraries, and file systems to accelerate development and reuse.
- **Safe On-Orbit Software Updates**  
Allow fault-tolerant update strategies (e.g., ping-pong schemes) with minimal mission risk.



Open Standard ISA



Open-Source Implementations


# Radiation Effect on SRAM-based FPGAs

- **Single Event Effects (SEEs)**

Energetic particles (e.g., protons, heavy ions) can strike sensitive regions in a transistor, causing transient or permanent effects.

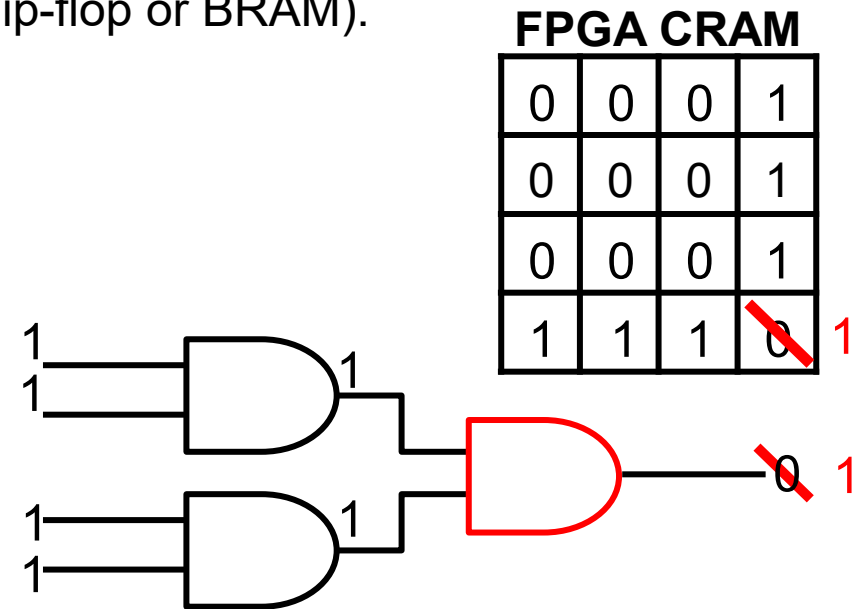
- **Single Event Upset (SEU)**

If the collected charge is sufficient to change the state of a logic element (e.g., SRAM cell), it results in a soft error.



- **SEUs** flip bits in the FPGA configuration memory, potentially altering the digital design and system functionality:


- **LUTs** – Changes logic functionality.
- **Routing** – Connect, disconnect or short wires.
- **Memory** – Corrupt the memory state (such as in a flip-flop or BRAM).



# Fault Tolerance for Soft RISC-V Systems


## Passive Fault Tolerance with Redundancy

- **Mask:** Continue operation by hiding faults using redundancy or correction logic.
- **Distributed Triple Modular Redundancy (dTMR)**  
Replicates all logic three times and inserts triplicated majority voters between modules to ensure correct outputs when one copy fails.



## Active Fault Tolerance during Operation

- **Repair:** Correct faults in hardware configuration or memory.
- **Configuration Scrubbing**  
Periodically reads and corrects upsets in the FPGA's configuration memory using a known-good bitstream or parity data.
- **Flip-Flop Resynchronization**  
Recovers from transient faults in sequential logic by reloading or re-aligning state



# Fault Tolerant RISC-V SoCs in FPGAs

## Problem: Redundancy Alone Is Not Always Enough for High Reliable Systems

- **Redundancy Benefits:** Can mask faults without disrupting system operation. Studies show reliability improvements ranging from **3×** [Olive et al., 2020] to **27×** [Keller et al., 2017].
- **Effectiveness Varies:** Depends on FPGA architecture, application type, and implementation strategy [Bolchini et al., 2006].
- **Synchronization Challenge:** Long-term protection requires failed modules to be resynchronized with healthy ones. Without this, redundancy degrades over time [Johnson, 2010].
- **Toolchain Limitations:** Many automated tools operate only at the functional level, ignoring physical layout constraints critical for fault tolerance [Cannon et al., 2019].
- **Design Implication:** Effective mitigation demands **fault injection**, **radiation testing**, and **systematic fault analysis** to ensure real-world reliability.

# Dissertation Contributions

**Goal:** Create Highest Reliable RISC-V Soft Processors on SRAM FPGAs

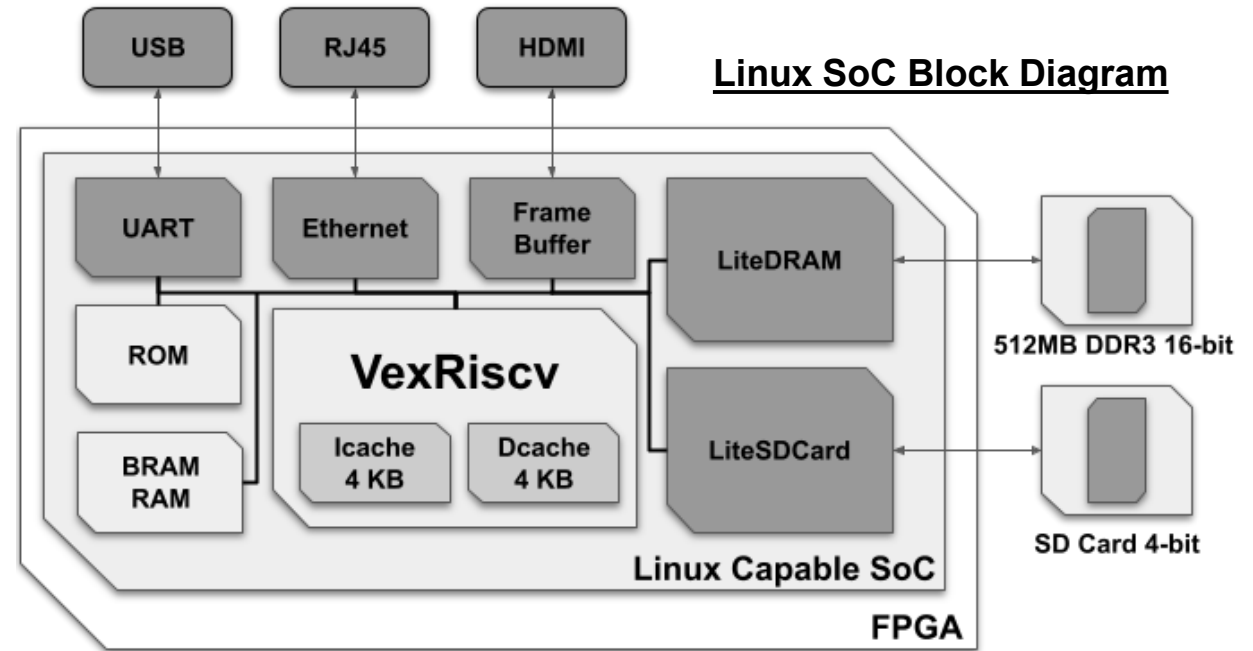

- Improved fault tolerance for RISC-V Linux SoCs on SRAM FPGAs
- Developed and applied structured fault analysis across multiple fault tolerant systems
  - Assisted by Bitstream Fault Analysis Tool (BFAT)
- Exposed TMR limitations and refined additional mitigation strategies
  - Verified 19.8× MTBF improvement for ECC implementation of open source DDR4 controller
  - Achieved 22.8× MTBF improvement with SEU-Aware placement mitigation
- Demonstrated failure analysis and identification for 98% of observed radiation-induced failures

# Designs Under Test

## Baseline Bare Metal RISC-V Processor

- Open-source Scala VexRiscv Processor
- 4KB unified instruction and data memory
- Serial output for basic telemetry

Baseline Block Diagram




## Linux RISC-V System on Chip (SoC)

- Expanded memory with external DDR3 SRAM.
- Processor support for Instruction and Data cache with memory management to support address virtualization.
- Advanced interfaces that support different boot options.
- Buildroot embedded Linux OS kernel
- LinuxOnLitex to rapid deploy FPGA implementations




# FPGA Architecture and Development Boards

- AMD-Xilinx FPGA 7 series (28nm) and UltraScale (20nm) Architectures
- The Processor SoCs utilized the DDR3/4, SD Card, and Ethernet.
- Different development boards for different variations over the period of the work



Alpha Data SDEV Kit  
AMD-Xilinx UltraScale  
Bare Metal RISC-V



Nexys Video  
AMD-Xilinx Series 7  
Linux 16-bit DDR SoC





Antmicro DDR Tester  
AMD-Xilinx Series 7  
Linux 64-bit DDR SoC



# Fault Injection for CRAM SEU Emulation


- **Purpose:** Evaluate system vulnerability by injecting controlled errors
- **Technique:** Emulate SEUs by flipping CRAM bits via JTAG commands
- **Target:** Logic & routing configuration bits (flip-flops excluded via glut mask)
- **Method:** Compare faulty vs. golden output to detect failure modes
- **Sampling:** Random fault locations provide a statistical reliability profile
- **Failure Sensitivity Formula:**

$$\text{Failure Sensitivity} = \frac{\text{Number of Failures}}{\text{Number of Faults Injected}}$$




# Radiation Testing

- **Purpose:** Radiation testing exposes FPGA designs to a controlled beam to simulate faults caused by space-like radiation environments.
- **Key Parameters:**
  - **Flux** ( $\varphi$ ): Rate of radiation particles per unit area (particles/cm<sup>2</sup>/s).
  - **Fluence** ( $F$ ): Total particle exposure over time, calculated as  $F = \varphi \cdot t$ .
- **Sensitivity Metric:** Failure sensitivity is measured via:
  - **Cross-section:** Failures per unit flux ( $r = k/\varphi$ ).
  - **Mean Fluence to Failure:** Average exposure required to cause a fault.
- **Use Case:** Validates resilience of soft processors to real radiation-induced faults.
- **Challenges:**
  - Limited access to calibrated radiation sources and high costs restrict test scope.
  - High flux rates can overwhelm time-based or single-event-oriented mitigation techniques (e.g., scrubbing or TMR), leading to accumulated faults beyond system tolerance.





LANSCE Facility (Los Alamos, NM)



ChiPlr Facility (London, UK)

# Fault Injection and Radiation Testing

- **JCM Utility:** Logs CRAM upsets and performs scrubbing using a golden bitstream
- **Soft Processor Failure:** Detected via output mismatch or benchmark timeout against golden reference
- **RISC-V SoC Sensitivity:**
  - Uses **2.55× more FPGA resources**, much greater failure susceptibility
- **Root Causes of Increased Sensitivity:**
  - Higher I/O complexity
  - Larger logic footprint
  - Longer critical paths with more routing




Design Under Test	Utilization (LUTs)	Design Sensitivity	Cross-Section ( $cm$ )
Bare Metal RISC-V	2,665 (1.96%)	0.05%	$5.14 \times 10^{-10}$
Linux RISC-V SoC	6,791 ( 5.0%)	0.774%	$2.09 \times 10^{-9}$
Ratio	2.55×	15.5×	4.00×




# Fault Tolerance through TMR and Scrubbing

- **Tool Purpose:** SpyDrNet enables netlist parsing and transformation for FPGA designs.
- **TMR Insertion:** Supports automated distributed TMR (dTMR / fine-grain TMR) at the netlist level.
- **Flip-Flop Feedback:** Injects feedback loops into flip-flops to support next-cycle resynchronization.
- **Customizability:** Designers can target specific modules, logic types, or hierarchy levels.
- **Scalability:** TMR is applied programmatically, avoiding manual duplication and wiring.

TMR Utilization Floorplan



SpyDrNet Tool Flow



TMR Utilization Table

Design	LUT	LUTRAM	FF	BRAM
16-bit non-TMR	6,791 ( 5.0%)	319 ( 0.7%)	5,506 ( 2.1%)	43.0 (11.8%)
16-bit TMR	27,916 (20.9%)	957 ( 2.1%)	16,512 ( 6.2%)	129.0 (35.3%)
Cost Ratio	4.11×	3.00×	3.00×	3.00×

# TMR Fault Tolerance Results

## Fault Injection Results

- **Unmitigated Sensitivity:** 0.774%
  - Baseline reliability (1.00×)
- **TMR Sensitivity:** 0.055%
  - **14.07× improvement**

Design	DDR Type	Injections	Failures	Sensitivity	CV	Improvement
16-bit	Unmitigated	41,206	319	0.774%	0.056	1.00×
	TMR	504,258	280	0.055%	0.060	14.07×

## Radiation Testing @ LANSCE 2021

- **Unmitigated**
  - 76 failures @  $3.64 \times 10^{10}$  n/cm<sup>2</sup>
  - Cross Section:  **$2.09 \times 10^{-9}$  cm<sup>2</sup>**
- **TMR Mitigated**
  - 27 failures @  $1.91 \times 10^{11}$  n/cm<sup>2</sup>
  - Cross Section:  **$1.42 \times 10^{-10}$  cm<sup>2</sup>**
- **MTBF Gain: 14.7× improvement**
- **What is MTBF?**
  - **Mean Time Between Failures (MTBF)** estimates the average time a system operates before failing.
  - Higher MTBF means **greater reliability**—TMR allows the system to run **~14.7× longer** on average without failure.

Design	Fluence (n/cm <sup>2</sup> )	Observed CRAM Upsets	Failures	Cross Section (cm <sup>2</sup> )	+95% Confidence -95% Confidence	Reduction
Linux-VexRiscv Unmitigated	$3.64 \times 10^{10}$	11908	76	$2.09 \times 10^{-9}$	$2.57 \times 10^{-9}$ $1.61 \times 10^{-9}$	1×
Linux-VexRiscv TMR	$1.91 \times 10^{11}$	59014	27	$1.42 \times 10^{-10}$	$3.47 \times 10^{-10}$ $4.82 \times 10^{-11}$	14.72×

# Fault Analysis Methods

- **Tile Analysis**

- Maps failure locations using frame address and logical tile mapping
- Highlights critical areas (e.g., DDR I/O, clock trees, unmitigated logic)

- **BFAT (Bitstream Fault Analysis Tool)**

A suite of Python-based static analysis tools applied to post-implementation DCPs:

- **1. FPGA Resource Analysis**

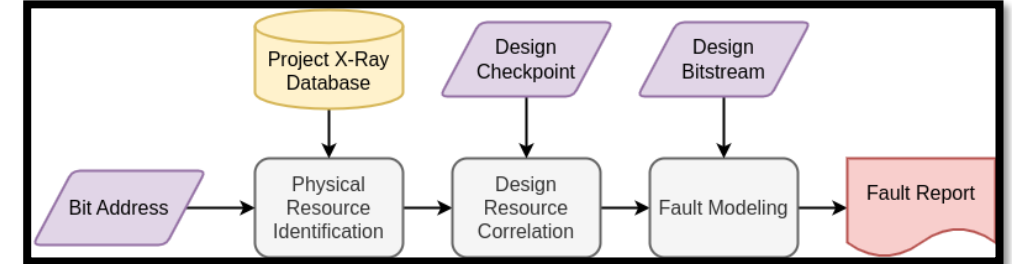
- Classifies failures by resource type: interconnect, I/O, clocking, CLBs

- **2. Design Function Analysis**

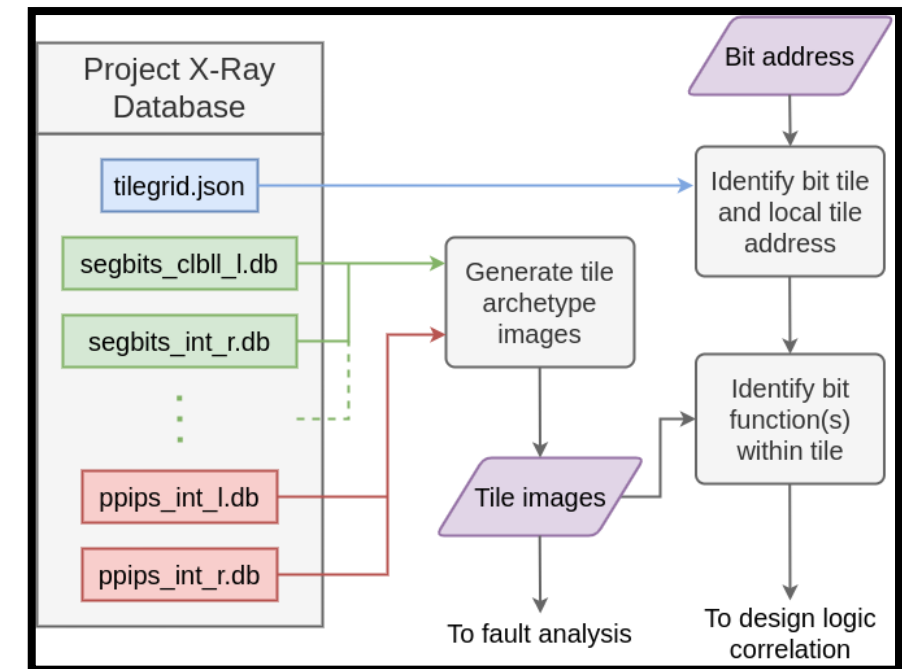
- Maps failures to logical modules (e.g., CPU, DDR controller, I/O)

- **3. Static Net Sensitivity Analysis**

- Identifies vulnerable, untriplcated PIPs and BELs
    - Excludes TMR, clock, and global nets




**BFAT Flow**





# TMR Linux RISC-V 16-bit DDR BFAT Report



## Tile-Based Fault & Resource Analysis (FPGA 2023)

- **Red** = radiation failures, **Green** = injection failures
- Most failures occur in **DDR I/O region** and **clock network tiles**
- **Over 70%** of failures caused by **interconnect faults**
- **Over 80%** tied to **DDR interface activity**
- TMR tools can't protect most interconnect faults related to the I/O primitives
- Next Slide: Example of a mapped fault in the implemented design related to the DDR interface

Resource	Occurrences (radiation test)	Occurrences (random fault injection)	Occurrences (total)	Estimated Bits
Interconnect	13 (76.4%)	75 (70.8%)	88 (71.5%)	17,855
I/O Pin	3 (17.6%)	20 (18.9%)	23 (18.7%)	4,667
Clocking	1 (5.9%)	6 (5.7%)	7 (5.7%)	1,420
CLB	0 (0.0%)	4 (3.8%)	4 (3.3%)	812
Undefined	0 (0.0%)	1 (0.9%)	1 (0.8%)	203
Total	17	106	123	24,956 (0.042%)*

\* Percentage is in comparison to the 59,145,600 type 0 CRAM locations in the XC7A200T

Table 6.2: FPGA Device Resources Causing Design Failure.

# DDR Input Upset in Routing

Bit Group 1

Failure Bits:

bit\_00402785\_100\_15 (0->1)

INT\_R\_X79Y149 - SW6BEG2 2-20 Routing Mux - Column Bit

Resource Design Name: INT\_R\_X79Y149/SW6BEG2

Shorts formed between net(s): ISERDESE2\_15\_n\_4 (initially connected), VexRiscvLiteSmpCluster\_Cc1\_Iw32Is4096Iy1\_Dw32Ds4096Dy1\_ITs4DTs4\_Ldw128\_Ood/dBridge

Affected PIPs:

WW2END2->>SW6BEG2 (activated)

Affected Resources:

VexRiscvLiteSmpCluster\_Cc1\_Iw32Is4096Iy1\_Dw32Ds4096Dy1\_ITs4DTs4\_Ldw128\_Ood/dBridge\_logic/io\_output\_rdata\_fifo/storage\_12\_reg\_10\_i\_12\_TMR\_0

VexRiscvLiteSmpCluster\_Cc1\_Iw32Is4096Iy1\_Dw32Ds4096Dy1\_ITs4DTs4\_Ldw128\_Ood/dBridge\_logic/io\_output\_rdata\_fifo/storage\_12\_reg\_8\_i\_10\_TMR\_0



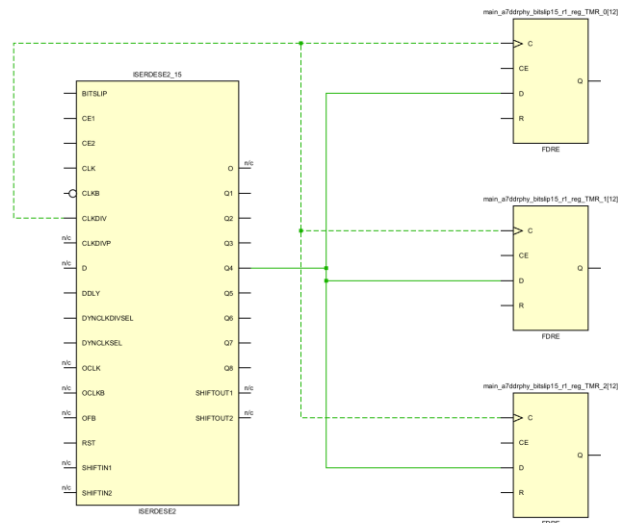
VexRiscvLiteSmpCluster\_Cc1\_Iw32Is4096Iy1\_Dw32Ds4096Dy1\_ITs4DTs4\_Ldw128\_Ood/dBridge\_logic/io\_output\_rdata\_fifo/storage\_12\_reg\_8\_i\_10\_TMR\_1

main\_a7ddrphy\_bitslip15\_r1\_reg\_TMR\_0[12]

main\_a7ddrphy\_bitslip15\_r1\_reg\_TMR\_1[12]

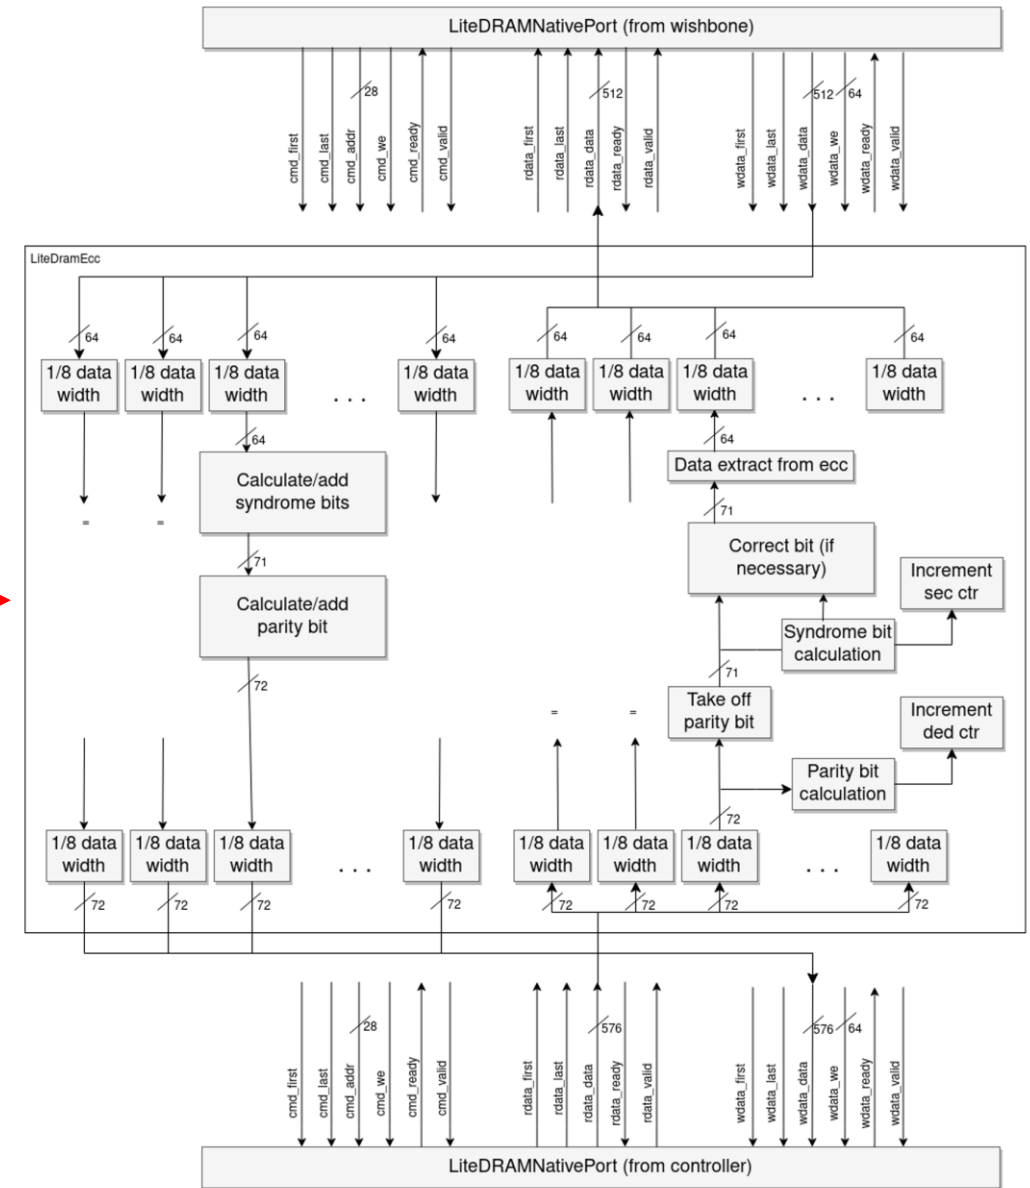
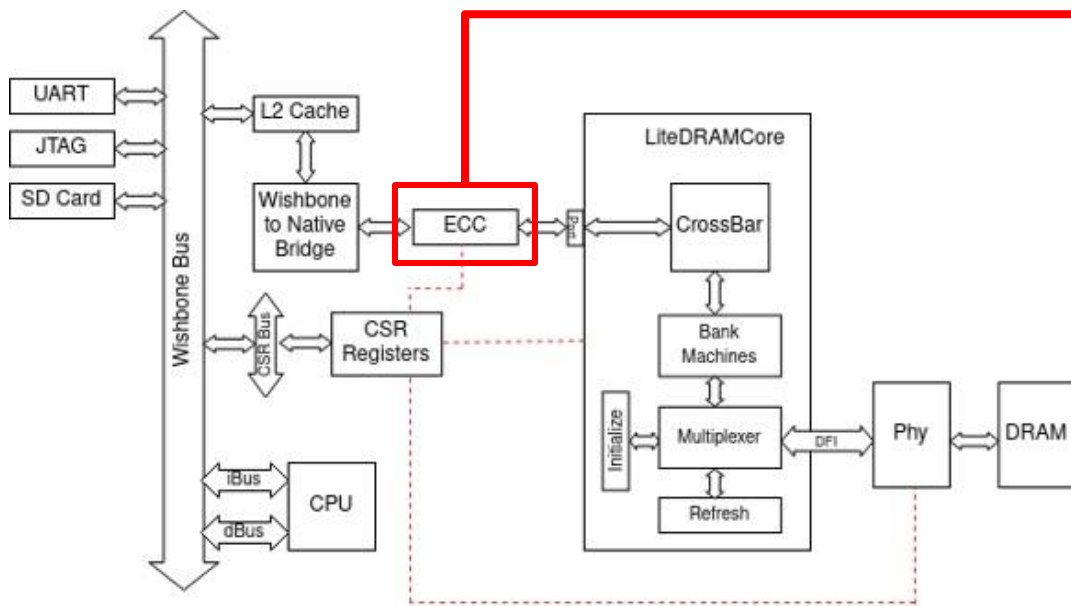
main\_a7ddrphy\_bitslip15\_r1\_reg\_TMR\_1[4]

main\_a7ddrphy\_bitslip15\_r1\_reg\_TMR\_2[12]



# ECC for TMR 64-bit SoC

- **Enable ECC Generator:** Modify LiteDRAM core to include Hamming (or SECDED) logic for data writes.
- **Integrate ECC Checker:** Add error detection and correction on data reads before passing to SoC/memory bus.
- **Extend Data Width:** Expand memory interface to accommodate ECC bits (e.g., 64b data + 8b ECC).





# TMR and ECC Fault Tolerance Results

## Fault Injection Results

- TMR 64-bit DDR Sensitivity: 0.230%
  - (Baseline TMR improvement : 5.55×)
- TMR & ECC Sensitivity: 0.0035%,
  - 35.06× improvement (~7× over just TMR)

Design	Injections	Failures	Sensitivity	CV	Improvement
Unmitigated	13,183	168	1.274%	0.077	1.0×
ECC	81,299	886	1.090%	0.033	1.17×
TMR	54,835	126	0.230%	0.089	5.55×
TMR & ECC	220,687	78	0.035%	0.113	36.06×

## Radiation Testing @ ChiPlr 2023

- Unmitigated:
  - 161 failures @  $4.21 \times 10^{10}$  n/cm<sup>2</sup>
  - Cross Section:  $3.82 \times 10^{-9}$  cm<sup>2</sup>
- TMR-Protected:
  - 105 failures @  $5.45 \times 10^{11}$  n/cm<sup>2</sup>
  - Cross Section:  $1.93 \times 10^{-10}$  cm<sup>2</sup>
- MTBF Gain: 19.84×

Design	Fluence (n/cm <sup>2</sup> )	Observed Cram Upsets	Failures	Cross Section (cm <sup>2</sup> )	+95% Confidence -95% Confidence	Reduction
64-bit DDR SoC Unmitigated	$4.21 \times 10^{10}$	15656	161	$3.82 \times 10^{-9}$	$4.42 \times 10^{-9}$ $3.22 \times 10^{-9}$	1×
72-bit DDR SoC TMR & ECC	$5.45 \times 10^{11}$	198225	105	$1.93 \times 10^{-10}$	$2.30 \times 10^{-10}$ $1.55 \times 10^{-10}$	19.84×

## BFAT Failure Analysis

- Interconnect single-point failures reduced by 90%

Resource	Occurrences (TMR)	Occurrences (TMR & ECC)	Estimated Bits (TMR)	Estimated Bits (TMR & ECC)
Interconnect	95 (82.6%)	31 (48.4%)	75,630	6,941
I/O Pin	10 (8.7%)	18 (28.1%)	7,961	4,030
Clocking	0 (0.0%)	10 (15.6%)	0	2,239
CLB	9 (7.8%)	5 (7.8%)	7,164	1,120
Undefined	1 (0.9%)	0 (0.0%)	796	0
Total	115	64	91,552 (0.23%)*	14,330 (0.036%)*

\*Percentage is in comparison to the 39,805,312 type 0 Cram locations in the XC7K160T


# SEU-Aware Placement for TMR 16-bit SoC

## Striping for TMR Isolation

- Assigns each TMR domain to a vertical FPGA stripe.
- Reduces routing overlap and PIP/switchbox sharing.
- Enforced via suffix naming (`_TMR_0/1/2`) + regex constraints.
- Limits the sharing of VCC/GND primitives between TMR domains.
- Improves fault isolation, but less effective in dense SoCs with the usage of special resources like BRAMs and DSPs.

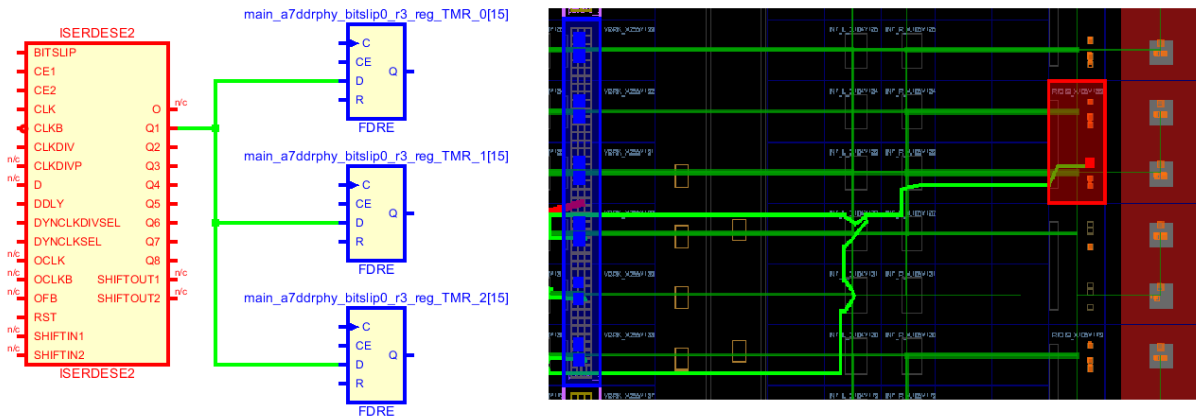
```
create_pblock pblock_dut_tmr_0
resize_pblock [get_pblocks pblock_dut_tmr_0] -add {
    SLICE_X0Y0:SLICE_X1Y249
    SLICE_X6Y0:SLICE_X7Y249
    ...
    SLICE_X96Y0:SLICE_X97Y249
    SLICE_X102Y0:SLICE_X103Y249
}
set_property IS_SOFT 0 [get_pblocks pblock_dut_tmr_0]
add_cells_to_pblock [get_pblocks pblock_dut_tmr_0] \
    [get_cells -hierarchical -regex .*TMR(_VOTER)?_0.*
    -filter IS_PRIMITIVE==1]
```

Example Striping TCL Script



# SEU-Aware Placement for TMR 16-bit SoC

## Input Path Example



```
create_pblock piserd_0
add_cells_to_pblock [get_pblocks piserd_0] [get_cells -quiet [list \
    main_a7ddrphy_bitslip0_r3_reg_TMR_0[15] \
    main_a7ddrphy_bitslip0_r3_reg_TMR_1[15] \
    main_a7ddrphy_bitslip0_r3_reg_TMR_2[15]]]
resize_pblock [get_pblocks piserd_0] -add {SLICE_X162Y180:SLICE_X163Y185}
```

## ISERDES TCL Example

## Output Path (OSERDES + LUT Voters)

- I/O primitives (e.g., OSERDES) can't be triplicated, expose critical vulnerabilities
- **LUT voters** reduce triplicated logic to a single value before OSERDES
- Default tool placement leads to long, SEU-prone nets
- **Mitigation:** Constrain voter LUTs near OSERDES blocks to reduce PIPs
- Result: 2.5× reduction in routing sensitivity (10.4 → 4.1 PIPs)

## Input Path (ISERDES + TMR Flip-Flops)

- ISERDES drives **triplicated flip-flops**, vulnerable if widely spread
- Default tool placement increases interconnect CRAM exposure
- **Mitigation:** Group flip-flops near each ISERDES block via pblocks
- Result: 3.2× reduction in PIPs (24.6 → 7.8 PIPs)



# Static Analysis on Unmitigated Routes

- **Objective:** Quantify and reduce single-point failure risks in configuration memory
- **Approach:** Apply incremental mitigations and re-analyze post-implementation design
- **Tool Used:** BFAT (Bitstream Fault Analysis Tool) for net-level sensitivity tracking
- **Scope:** Focused on unmitigated PIPs, BELs, and routing CRAM bits
- **Key Result:**
  - **Unmitigated Design:** 2.95% of CRAM bits sensitive
  - **Fully Mitigated Design:** 0.056% sensitive  
→ **37.5× reduction** in vulnerability

Design Variant	Description of Mitigations
TMR with One Clock	Baseline 16-bit TMR Linux SoC with a single shared clock tree. No placement constraints.
TMR Clocks	Triplicated BUFG primitives added, with constraints to preserve clock buffer separation.
TMR Clocks + Striping	Adds domain striping: physical separation of logic for each TMR domain via pblocks.
TMR Clocks + DDR I/O	TMR BUFGs and flip-flops around OSERDES/ISERDES blocks to reduce I/O net sensitivity.
TMR & All Mitigations	Final design incorporating all strategies: TMR clocks, domain striping, and placement-aware DDR I/O constraints.

Design	Unmitigated PIPs & BELs	Percent of CRAM*	Improvement
Unmitigated	1,238,982	2.95%	1.0×
TMR with One Clock	71,802	0.121%	17.26×
TMR Clocks	59,991	0.101%	20.65×
TMR Clocks + Striping	45,428	0.077%	27.27×
TMR Clocks + DDR I/O	31,631	0.053%	39.17×
TMR & All Mitigations	33,019	0.056%	37.52×

\*Calculated by dividing the number of unmitigated PIPs and BELs by the total number of Type 0 CRAM bits in the XC7A200T device (59,145,600).

# SEU-Aware PAR, TMR Fault Tolerance Results

## Fault Injection Results

- TMR 16-bit DDR Sensitivity: 0.039%
  - (Baseline TMR improvement : 14.74×)
- TMR & All Mitigations Sensitivity: 0.0019%
  - 31.02× improvement

Design Variant	Injections	Failures	Sensitivity (%)	CV	Improvement
Unmitigated	24,218	139	0.574%	0.085	1.00×
TMR with One Clock	603,366	235	0.039%	0.065	14.74×
TMR Clocks	121,162	37	0.031%	0.164	18.79×
TMR Clocks + Striping	155,635	45	0.029%	0.149	19.85×
TMR Clocks + DDR I/O	103,884	25	0.024%	0.200	23.85×
TMR & All Mitigations	270,209	50	0.019%	0.141	31.02×

## Radiation Testing @ LANSCE 2022

- Unmitigated:
  - 95 failures @  $3.67 \times 10^{10}$  n/cm<sup>2</sup>
  - Cross Section:  $2.59 \times 10^{-9}$  cm<sup>2</sup>
- TMR-Protected:
  - 37 failures @  $3.26 \times 10^{11}$  n/cm<sup>2</sup>
  - Cross Section:  $1.14 \times 10^{-10}$  cm<sup>2</sup>
- MTBF Gain: 22.78×

Design	Fluence (n/cm <sup>2</sup> )	Observed Cram Upsets	Failures	Cross Section (cm <sup>2</sup> )	+95% Confidence -95% Confidence	Reduction
Linux-VexRiscv (Unmitigated)	$3.67 \times 10^{10}$	16,239	95	$2.59 \times 10^{-9}$	$3.12 \times 10^{-9}$ $2.06 \times 10^{-9}$	1×
Linux-VexRiscv (TMR & All Mitigations)	$3.26 \times 10^{11}$	144,841	37	$1.14 \times 10^{-10}$	$2.70 \times 10^{-10}$ $3.38 \times 10^{-11}$	22.78×

## BFAT Failure Analysis

- Interconnect single-point failures reduced by 50%

Resource	Occurrences (original)	Occurrences (improved)	Estimated Bits (original)	Estimated Bits (improved)
Interconnect	88 (71.5%)	53 (57.0%)	17,855	8,762
I/O Pin	22 (17.9%)	20 (21.5%)	4,464	3,306
Clocking	7 (5.7%)	3 (3.2%)	1,420	496
CLB	4 (3.3%)	16 (17.2%)	812	2,646
Undefined	2 (1.6%)	1 (1.1%)	406	169
Total	123	93	24,956 (0.042%)*	15,379 (0.026%)*


\* Percentages are in comparison to the 59,145,600 type 0 Cram locations in the XC7A200T

# Fault Analysis Fault Tolerant Evaluation

- **Intrinsic device and design constraints** cap fault tolerance at  $\sim 50\times$  MTBF improvement
- **Residual failures** are largely due to vulnerabilities in the **DDR memory interface**
- **Unmitigatable single-point failures** include:
  - SEFIs (Single Event Functional Interrupts)
  - Critical clock routing mux failures

Failure Category	Count	Percent
Unmitigatable Failures	24.5	23.33%
XDC-Based Placement	12.0	11.43%
ADDR & CMD Parity	12.5	11.90%
DDR Error Accumulation	56.0	53.33%
Total	105	100%

**Table 9.1:** Failure Categories for the Last 5%



**Figure 9.1:** The Last 5% of Failures for the 72-bit TMR & ECC Design

# Potential Mitigations

- **ECC Data Scrubbing**

- Periodically reads and corrects memory upsets using error-correcting codes
- A scrub-to-upset ratio of **1:10** can reduce memory-related failures by **up to 65×**

- **Distributed Reset for I/O Primitives**

- Actively resets vulnerable I/O blocks (e.g., SERDES, IOBs) during faults
- Prevents persistent error states without full reconfiguration
- Especially useful for **non-triplicated I/O**, where redundancy cannot be applied

*[Cannon et al., “Improving the Reliability of TMR With Nontriplicated I/O on SRAM FPGAs,” 2019]*

- **DDR4 Address & Command Parity**

- Adds parity bits to control signals (e.g., address, command) for DDR4 memory
- Helps detect and isolate control path errors that could corrupt memory transactions

# Lessons Learned

- Majority of Fault Tolerant design failures are related to FPGA physical resource usage, placement, and routing.
  - Beyond synthesis-level TMR tools capabilities.
- TMR Fault tolerance is highly dependent on design choices and should always be coupled with fault injection and analysis.
- There is still opportunity for significant fault tolerance improvement at the cost of minimal overhead and design intrusion.
  - DDR ECC scrubbing would be a **2× MTBF** incremental improvement
  - Address & Command Parity would be a **1.5× MTBF** incremental improvement
  - Distributed Reset would be a **1.5× MTBF** incremental improvement

# Pushing Fault Tolerance Further

- Need SW & HDL functions beyond a simple benchmark.
  - (more realistic baseline).
- BRAM Scrubbing on ROMs and time-out on cached entries.
- Mitigation against configuration signals (i.e. xapp 1098, init\_b intercept)
- Dedicated Isolated Clk sources and internal clock routing.
  - (3x MMCMs with sync, isolated routes to BUFGs, redundant DDR clock source).
- Extended control for 72-bit DDR4 with byte module-level recovery.
  - via byte module reset and RAID 5 parity scheme.
- Hardened external TMR voters for critical reset lines.
- Kernel Modification for monitoring and checkpoint recovery.
  - Focus on optimizing uptime and reducing downtime.



# Publications and Presentations

- **Primary Author**
  - Neutron radiation testing of fault tolerant RISC-V soft processor on Xilinx SRAM-based FPGAs - 2019 IEEE Space Computing Conference (SCC)
  - Fault Injection Testing of Fault Tolerant RISC-V Soft Processors on Xilinx SRAM-based FPGAs – 2021 Journal of Radiation Effects Research and Engineering
  - Neutron radiation testing of a TMR VexRiscv soft processor on SRAM-based FPGAs – 2021 IEEE Transactions on Nuclear Science
  - Post-radiation Fault Analysis of a High Reliability FPGA Linux SoC – 2023 International Symposium on Field Programmable Gate Arrays
  - Neutron radiation testing of RISC-V TMR soft processors on SRAM-based FPGAs – 2023 IEEE Transactions on Nuclear Science
- **Secondary Author**
  - Post-Irradiation Fault Injection for Complex FPGA Designs – 2025 IEEE Transactions on Nuclear Science
- **Poster**
  - Neutron Radiation Results of Fault-Tolerant Soft RISC-V Linux SoCs on SRAM-based FPGAs – 2025 IEEE Nuclear and Space Radiation Effects Conference
- **Conference Presentations**
  - You've got to break a few FPGAs to go to Space! – 2022 Hackaday SuperCon
  - Evaluating a Cost-Effective, Open-Source, and Radiation-Tolerant FPGA SDR via Fault Injection for SmallSat Applications – 2025 SmallSat Salt Lake City

# Future Work

- Reaching the Fault Tolerance limit with new mitigation methods.
  - ECC Data scrubbing.
  - Distributed reset for IO primitives.
  - Address and Command Parity for DDR4.
- Exploring newer generations of FPGA and different vendors.
- Implement automation of placement constraints with tightly coupled fault injection and analysis.
- Look at different fault tolerant methods for external memory such as Reed-Solomon error correction codes or radiation intelligent memory Controllers.

# DOOM!



---

# Thank You

- Questions Please!